

GL5 Project

Requirements

Write a distributed application that you are passionate about. The application should fulfill a well defined task that's useful (not something that calculates the factorial of the number).

Having a UI is not strictly required, a working API is enough.

We require the application to rate limit user requests based on their IP address.

DevOps

- Observability:
 - Metrics:
 - The application/service should have some general purpose metrics (mem, cpu, request count) exported.
 - 1 metrics that allow us to track how many request we are processing
 - Label the metric in a way that allows you to get the success ratio (% of successful requests)
 - Label the metrics in a way that allows you to know many requests each http route processed.
 - Business logic metrics
 - Add a metrics that reflect a business value of your application, example:
 - Order count per item (for a shop website)
 - Friend request rejected (for a social media application)
 - Logs:
 - All logs emitted when processing a request should have the attribute `request_id` that uniquely identifies the request
 - For any given request you should either:
 - Generate a new request id (an UUID), or
 - Use whatever value found in the request header `X-Request-ID`
 - All logs should have the attribute `client_ip` which is the ip address of the caller
 - Trace:
 - Generate traces for your handlers
 - Add the same attributes you been asked to add in logs to the traces as well (`request_id` and `client_ip`)
- helm:
 - Write a helm chart for your service
 - Make these values configurable via the file `values.yaml`
 - `image.name`
 - `image.tag`

- http.port
- deployment.replicas
- Make it possible to add arbitrary `labels` to the deployment via the file values.yaml
- Make sure you define the correct env variables so that open telemetry tracing work as expected

Automation

- Automated infrastructure provisioning
- Automated initialization for workloads
- Multi-environment setup (Or at least the project should be capable of supporting extra environments)
- Well isolated and maintainable infrastructure layers (Microstacks)
- Automated Deployment

Deployment

- Usage of Microservices
 - Use all good practices of tracking workloads health and restoring them when needed
- Adequate choice of deployment strategy (+ explanation)
 - Note: Recreate and Rolling updates are not accepted
- Use of required configurations and secrets (treat those secrets carefully)
- Use an adequate network architecture and a well-defined set of network policies
- Have at least one endpoint exposed to the public

Shared

- Code quality
 - Well structured project tree
 - Good documentation, project stack overview + guides
- All work should be available in a git repository that we (Wassim and Meher) can access.
 - Project setup should be straightforward and predictable (we will build it on our machines), Use the adequate technologies to ensure that.
 - The repository should contain the
 - Application source code
 - The helm chart definition
 - Files needed to build the application
 - Documentation where the application architecture is well explained.
- We prefer each student to work individually, maximum is 3.
 - **The bigger the group the higher the expectations.**